

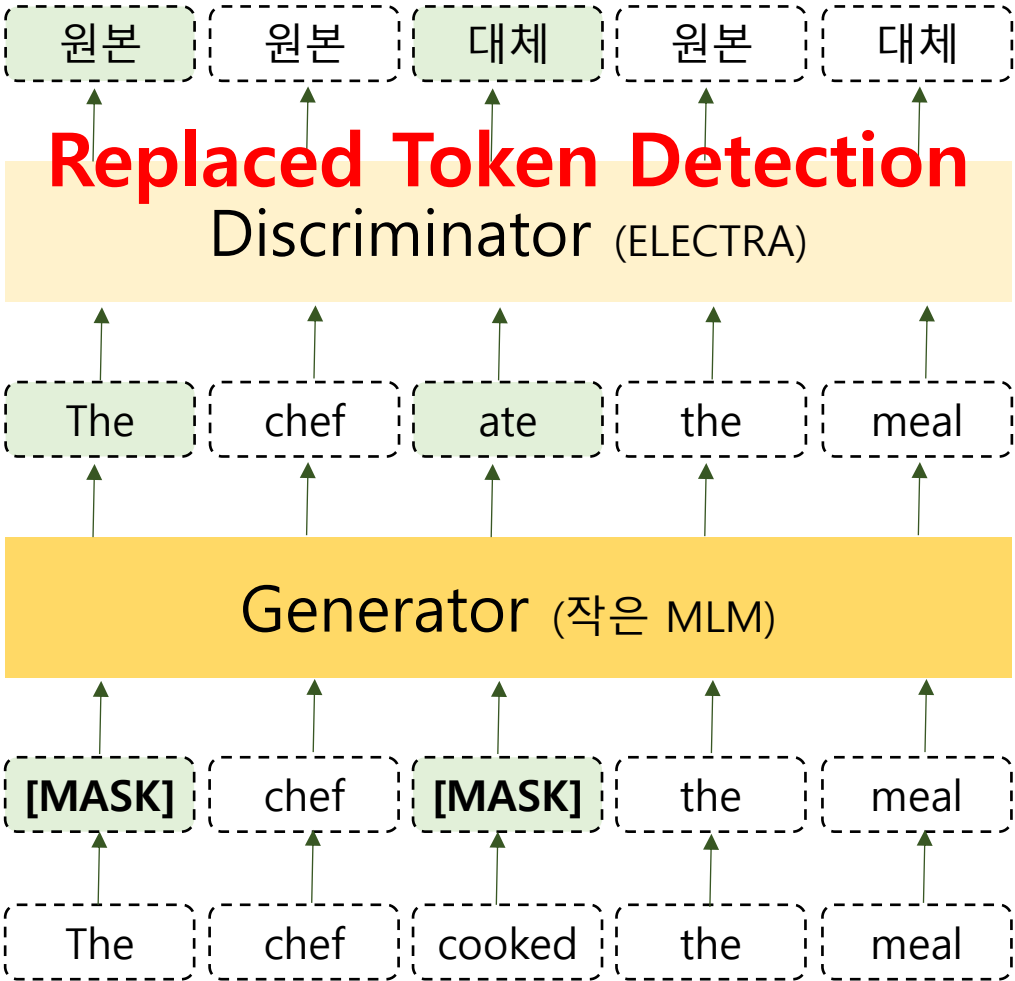
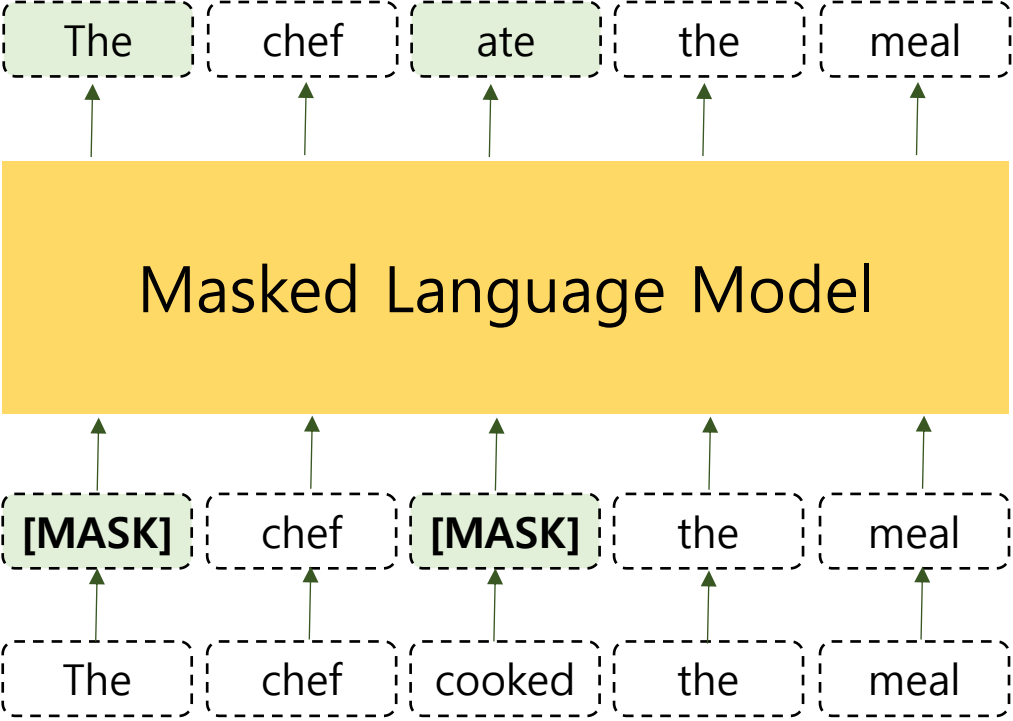
ELECTRA

10조
고현지 김병진 차지윤

Efficiently
Learning an
Encoder that
Classifies
Token
Replacements
Accurately

METHOD

BERT vs ELECTRA



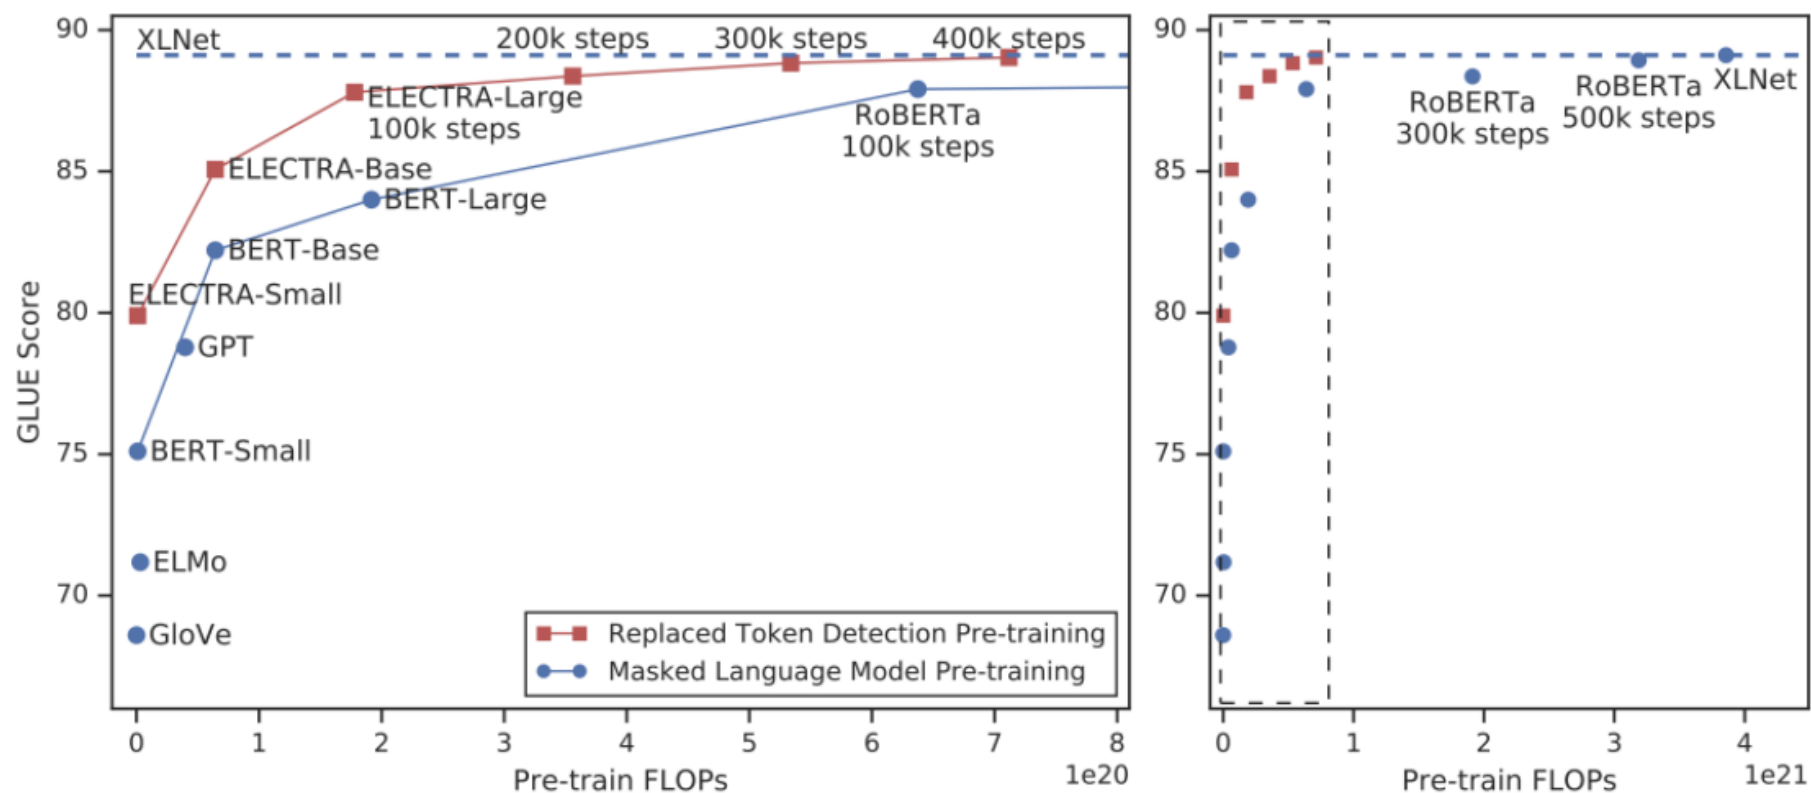
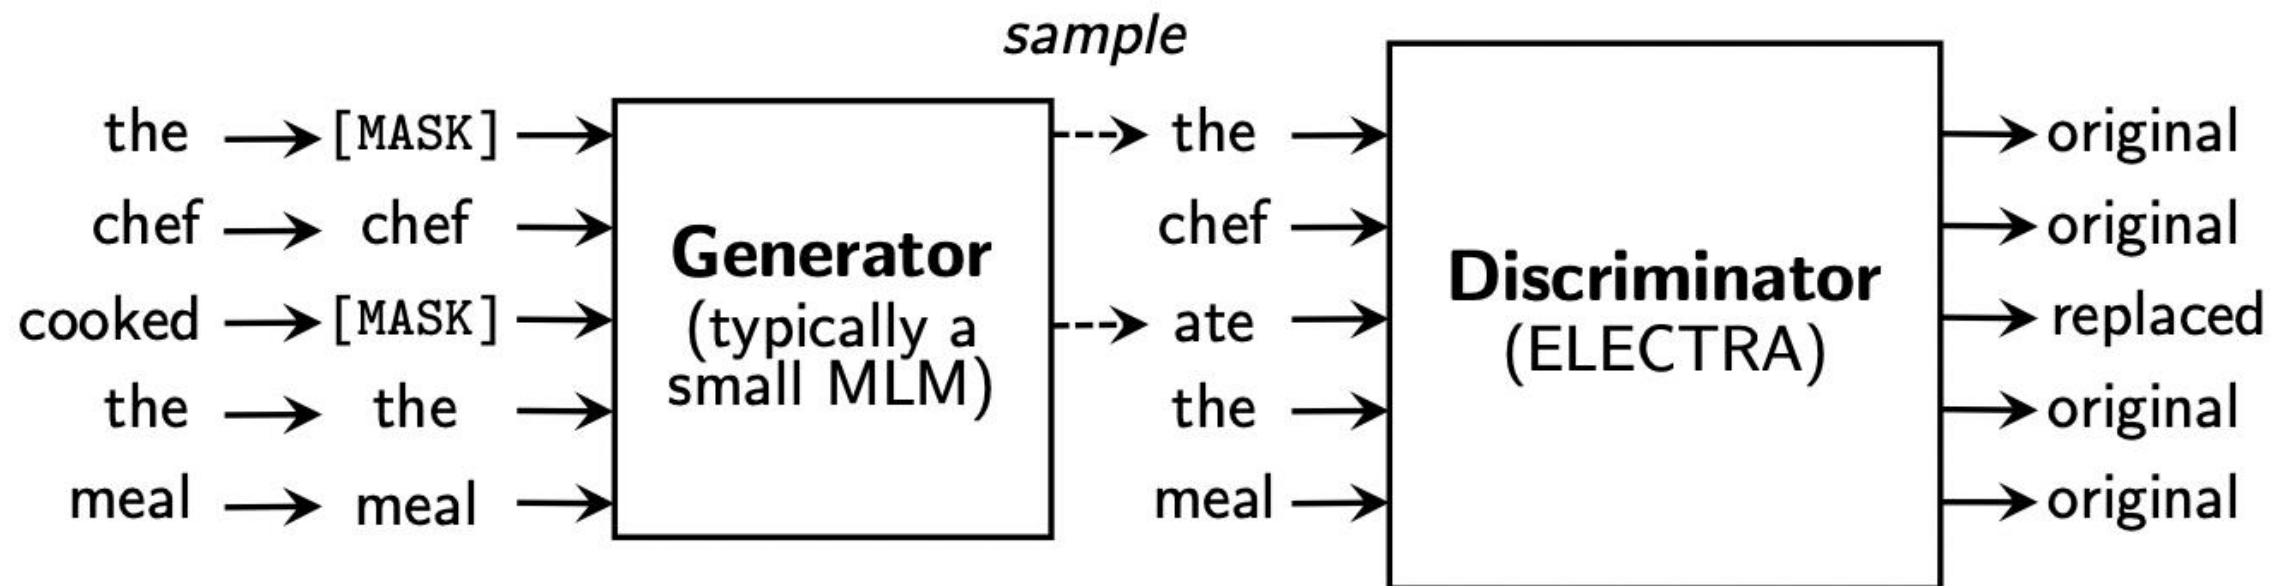
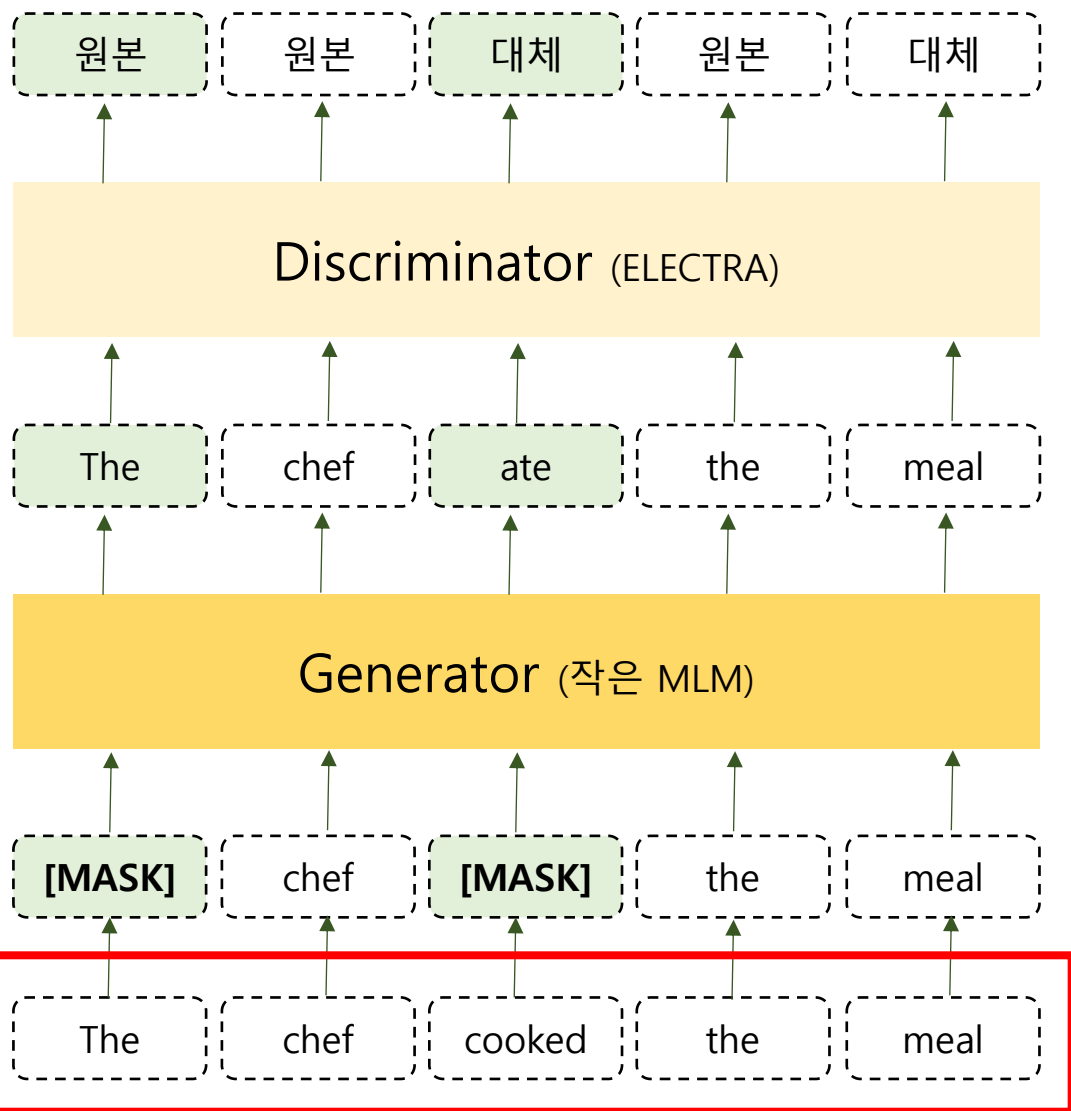


Figure 1: Replaced token detection pre-training consistently outperforms masked language model pre-training given the same compute budget. The left figure is a zoomed-in view of the dashed box.

METHOD



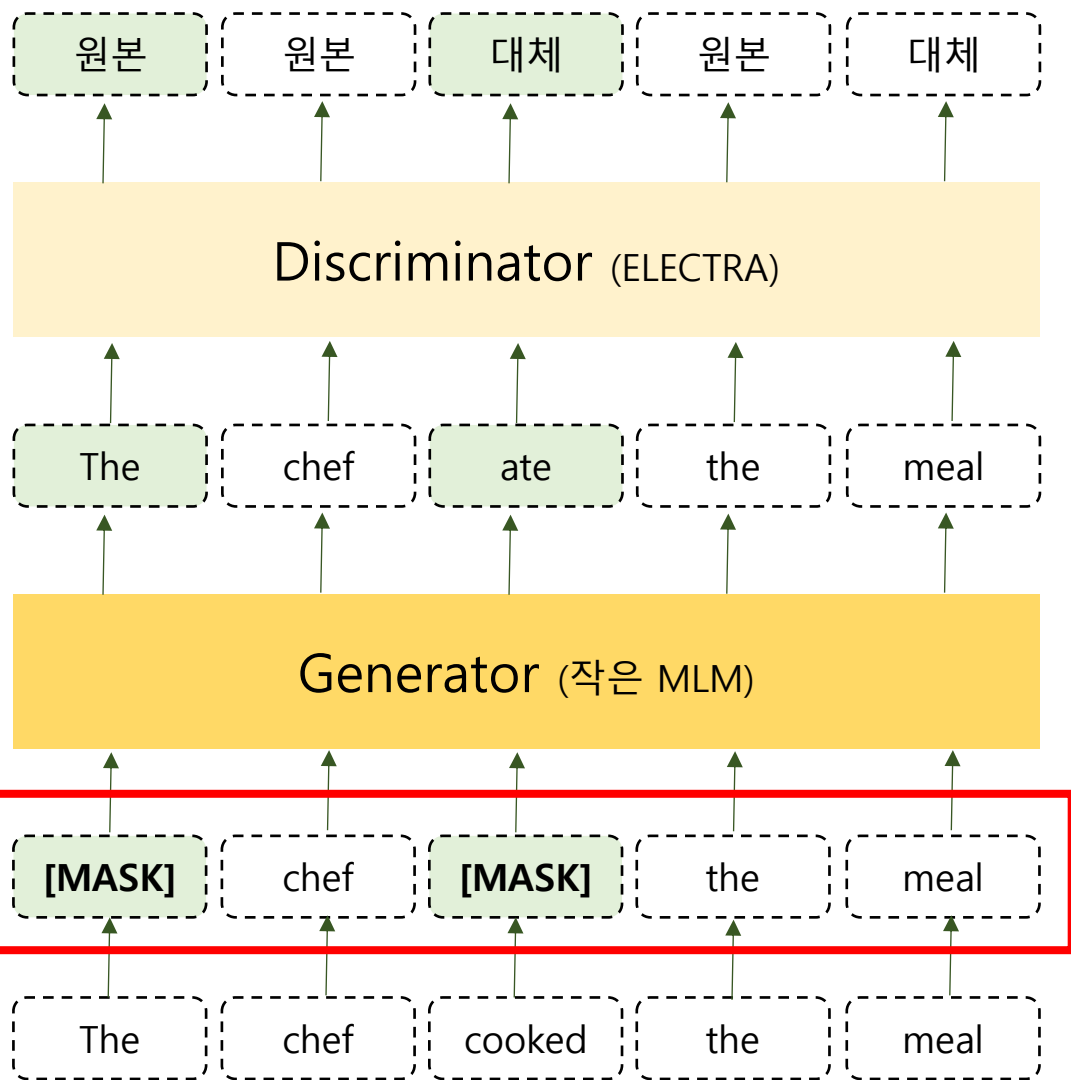
Generator G



1 $x = [x_1, x_2, \dots, x_n]$ 에 대해서
마스킹할 위치 $m = [m_1, m_2, \dots, m_k]$ 을 랜덤하게 결정

$$m_i \sim \text{unif}\{1, n\} \text{ for } i = 1 \text{ to } k$$

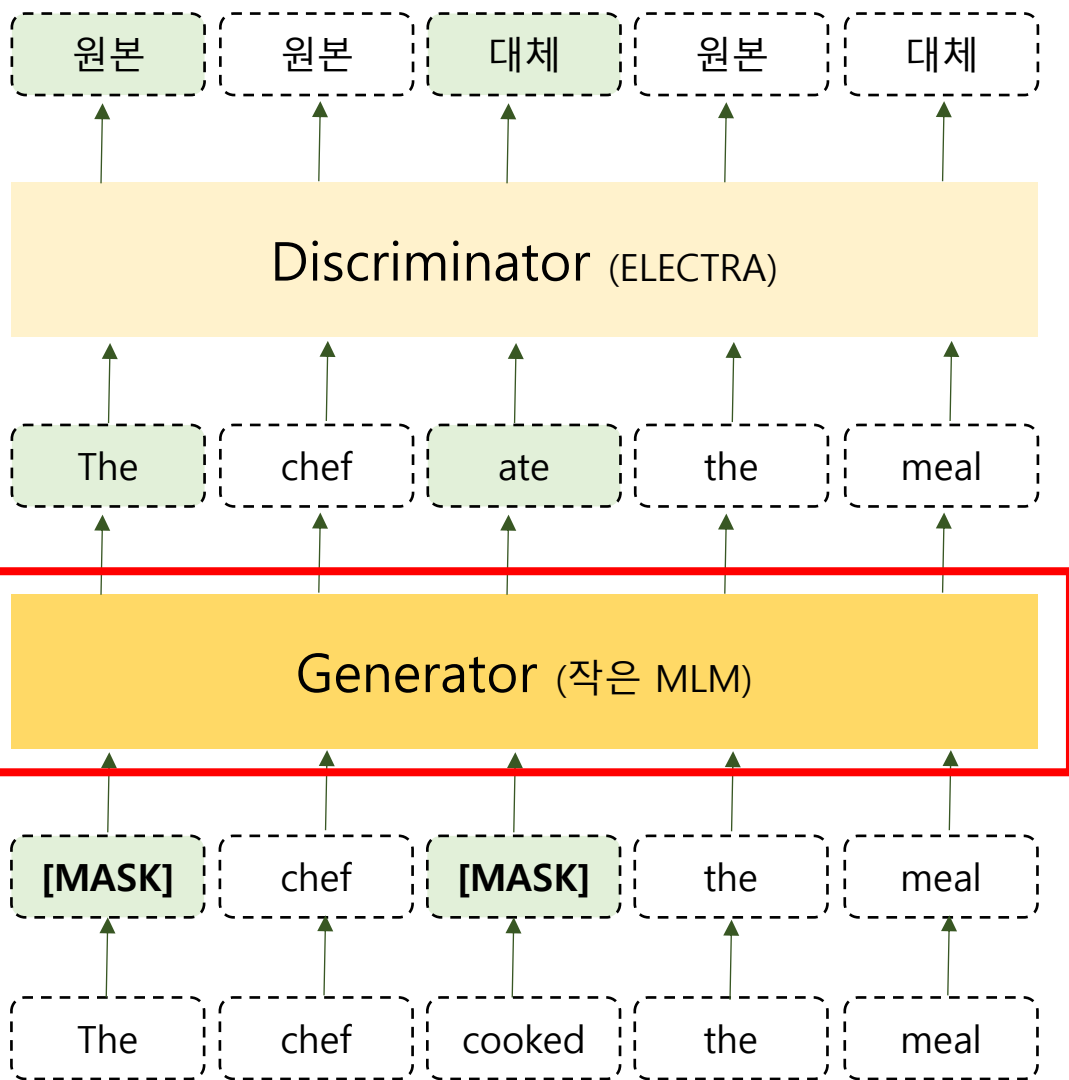
Generator G



2 마스킹된 토큰들은 [MASK] 토큰으로 대체

$$x^{masked} = REPLACE(x, m, [MASK])$$

Generator G



3 마스킹 된 입력 시퀀스 x^{masked} 에 대해서 generator는 아래와 같이 원래 토큰이 무엇인지 예측

$$p_G(x_t | x^{masked}) = \frac{\exp(e(x_t)^T h_G(x^{masked})_t)}{\sum_{x'} \exp(e(x')^T h_G(x^{masked})_t)}$$

$$p_G(x_t | x^{masked})$$

$\Rightarrow x^{masked}$ 라는 시퀀스 주어졌을 때 특정 위치 t에 단어 x_t 가 들어갈 확률

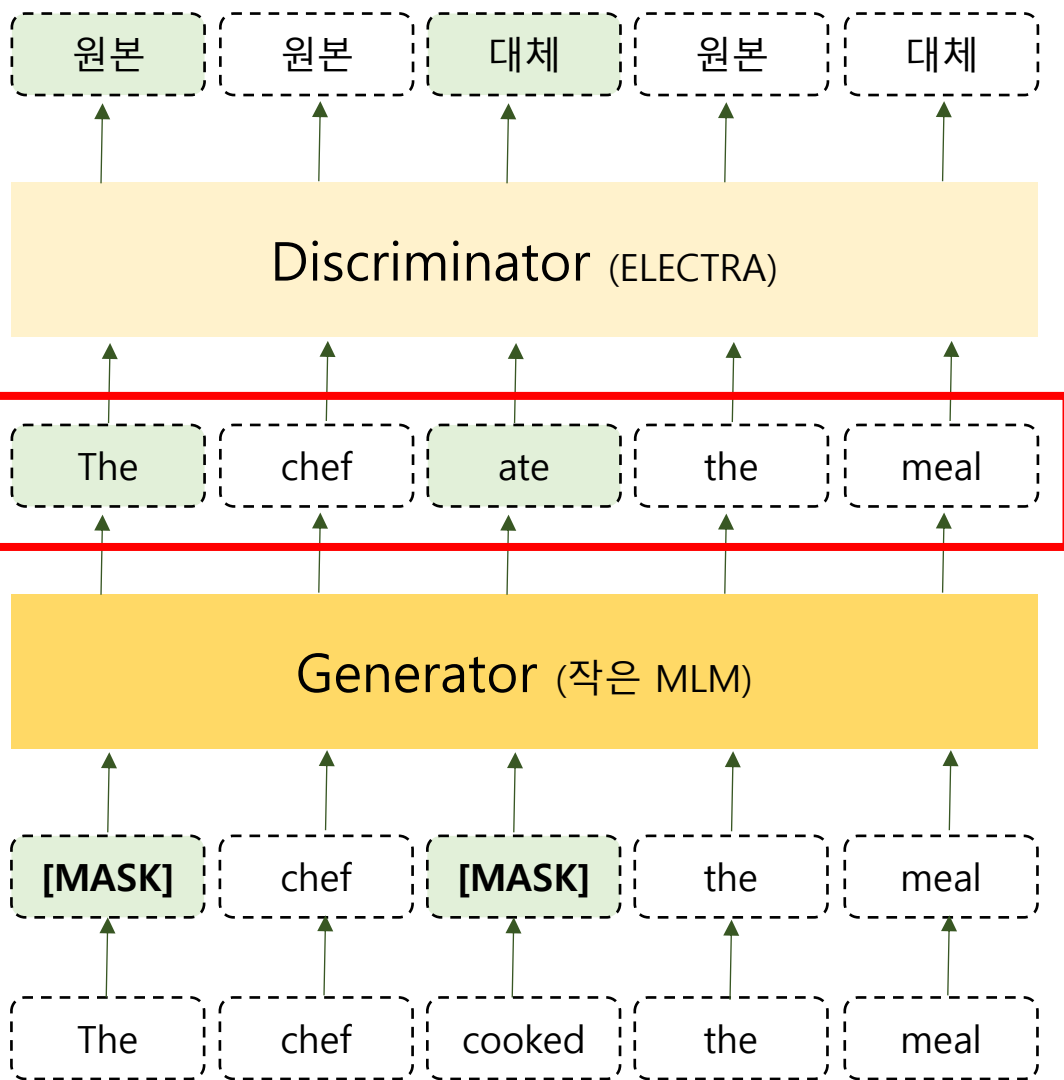
$$\frac{\exp(e \cdot x_i)}{\sum_j \exp(e \cdot x_j)}$$

\Rightarrow softmax 함수

$$e(x_t)^T h_G(x^{masked})_t$$

\Rightarrow token을 임베딩(e)하여 example의 hidden state 값을 곱

Discriminator D



1

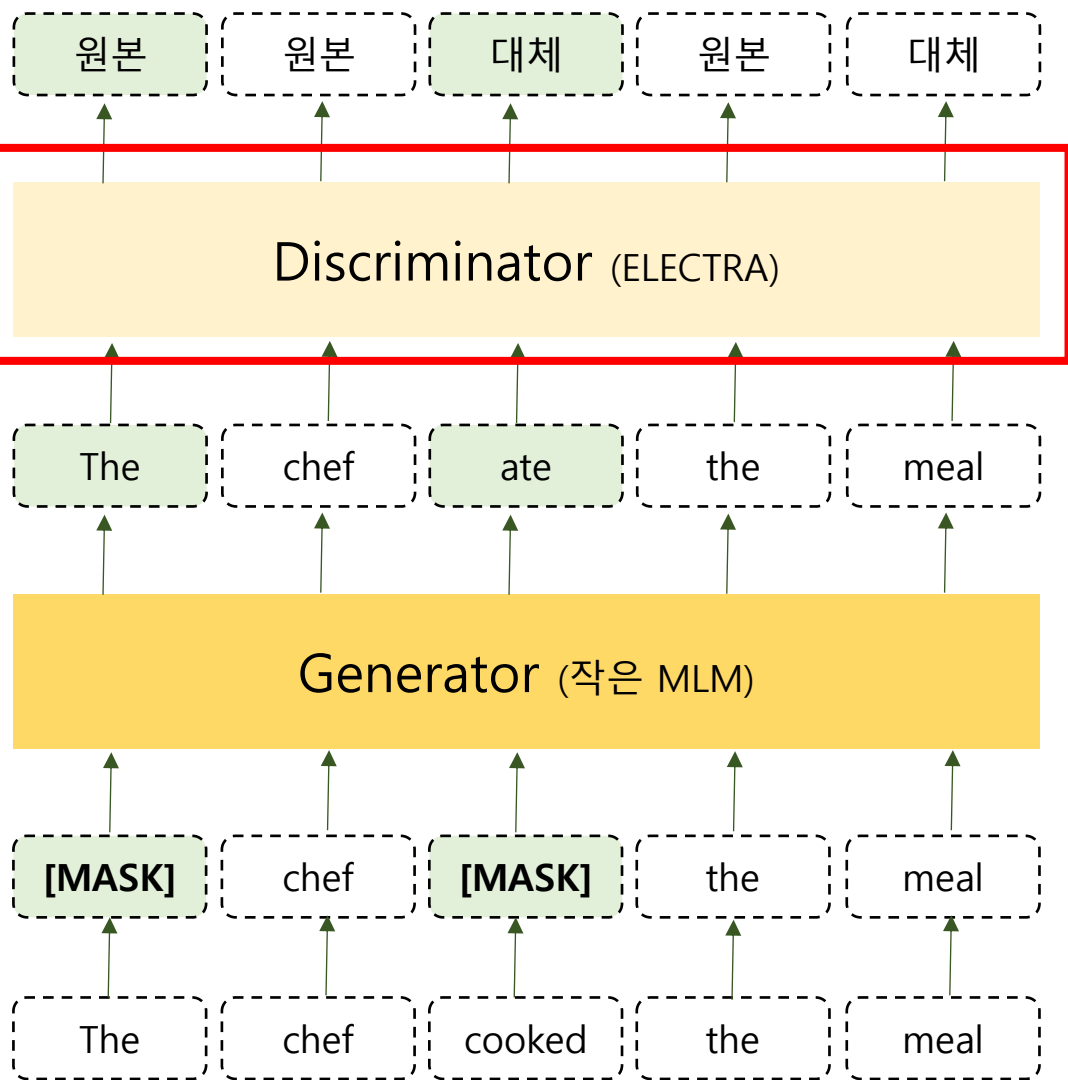
Discriminator D input 생성

[MASK]에서 $p_G(x_t|x)$ 으로 샘플링한 토큰으로 치환(corrupt)

$$x^{corrupt} = REPLACE(x, m, \hat{x})$$

$$\hat{x} \sim p_G(x_i|x^{masked}) \text{ for } i \in m$$

Discriminator D



2 Discriminator D는 $x^{corrupt}$ 가 원본 입력 토큰과 동일한건지, Generator G가 만들어낸 것인지 예측

$$D(x^{corrupt}, t) = \text{sigmoid}(w^T h_D(x^{corrupt})_t)$$

Target Class (이진)

- Original : 원본 문장의 토큰과 같은 토큰
- Replaced : 원본 문장의 토큰과 다른 토큰

Generator G와 Discriminator D 의 학습

Generator G

$$L_{MLM}(x, \Theta_G) = E(\sum_{i \in m} -\log p_G(x_i | x^{masked}))$$

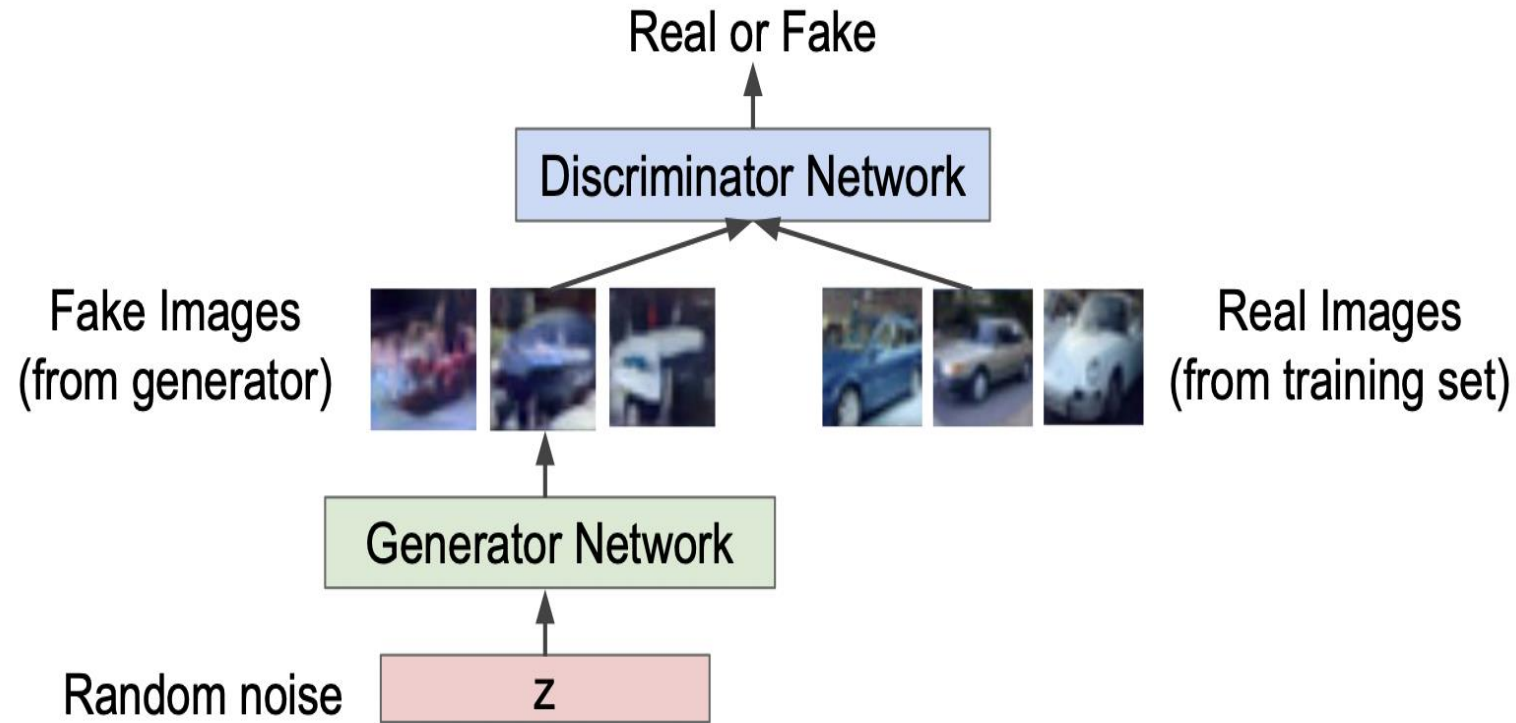
Discriminator D

$$L_{Disc}(x, \theta_D) = E(\sum_{t=1}^n -(x_t^{corrupt} = x_t) \log D(x^{corrupt}, t) - (x_t^{corrupt} \neq x_t) \log(1 - D(x^{corrupt}, t)))$$

대용량 코퍼스에 대해서 generator loss와 discriminator loss의 합을 최소화하도록 학습

$$\min_{\theta_G, \theta_D} \sum_{\mathbf{x} \in \mathcal{X}} \mathcal{L}_{MLM}(\mathbf{x}, \theta_G) + \lambda \mathcal{L}_{Disc}(\mathbf{x}, \theta_D) \quad \lambda = 50$$

GAN



Generator : 진짜같은 image를 만들어내서 Discriminator를 속임
Discriminator : Generator가 만들어낸 이미지인지 실제 이미지인지 구분

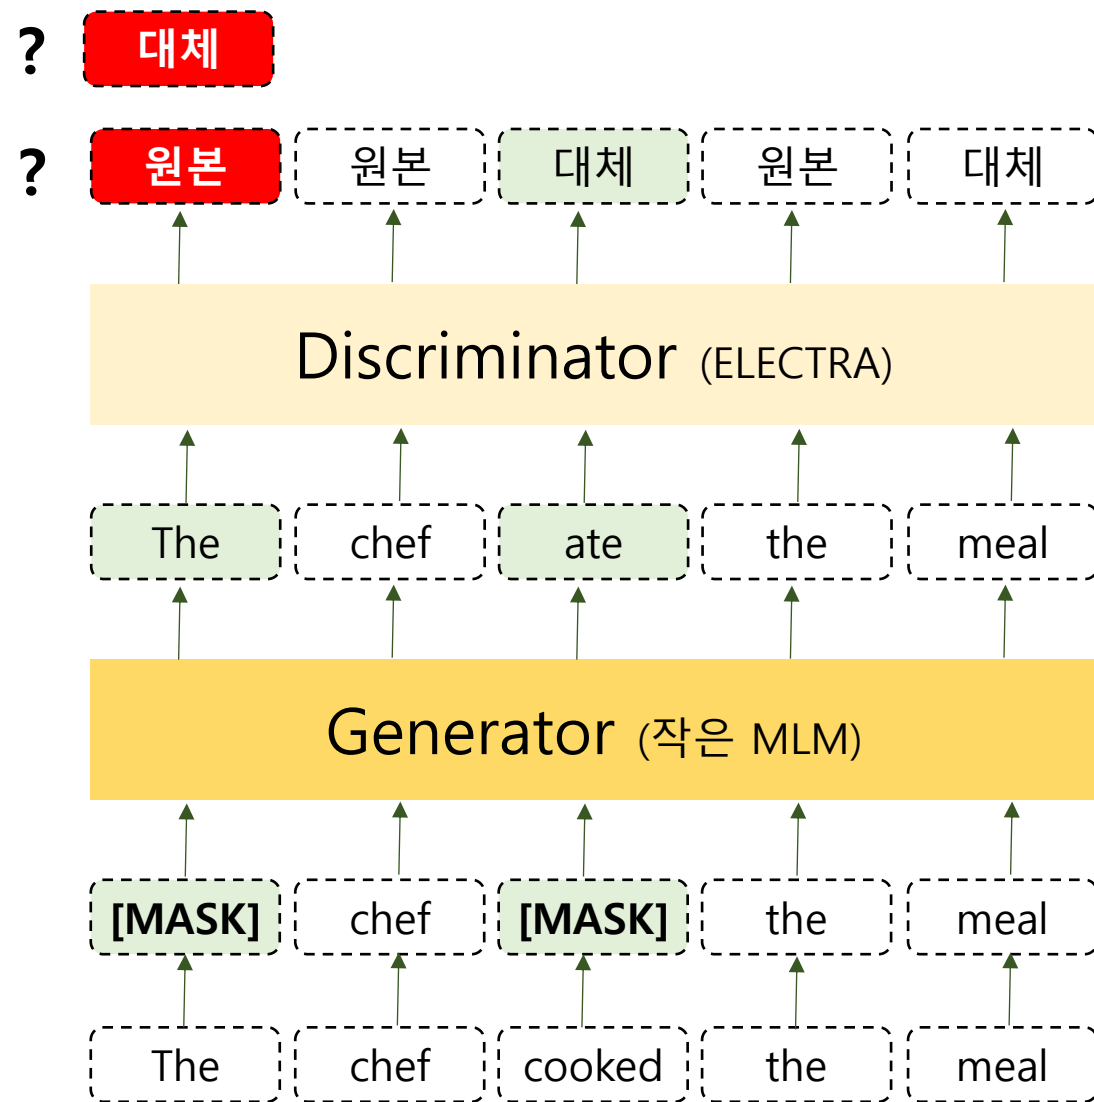
GAN vs ELECTRA

1. Generator가

원래 토큰과 동일한 토큰을 생성한다면?

GAN : negative sample (fake)로 간주

ELECTRA : positive sample로 간주



GAN vs ELECTRA

2. 학습 방법?

GAN : generator와 discriminator가 adversarial하게 학습

ELECTRA : generator와 discriminator가 각자 학습

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Alternate between:

1. **Gradient ascent** on discriminator

$$\max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

2. **Gradient descent** on generator

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$

GAN vs ELECTRA

3. Generator의 입력으로 노이즈 벡터?

GAN : 넣음
ELECTRA : 넣지 않음

EXPERIMENTS

Experiments

Experiments : SETUP


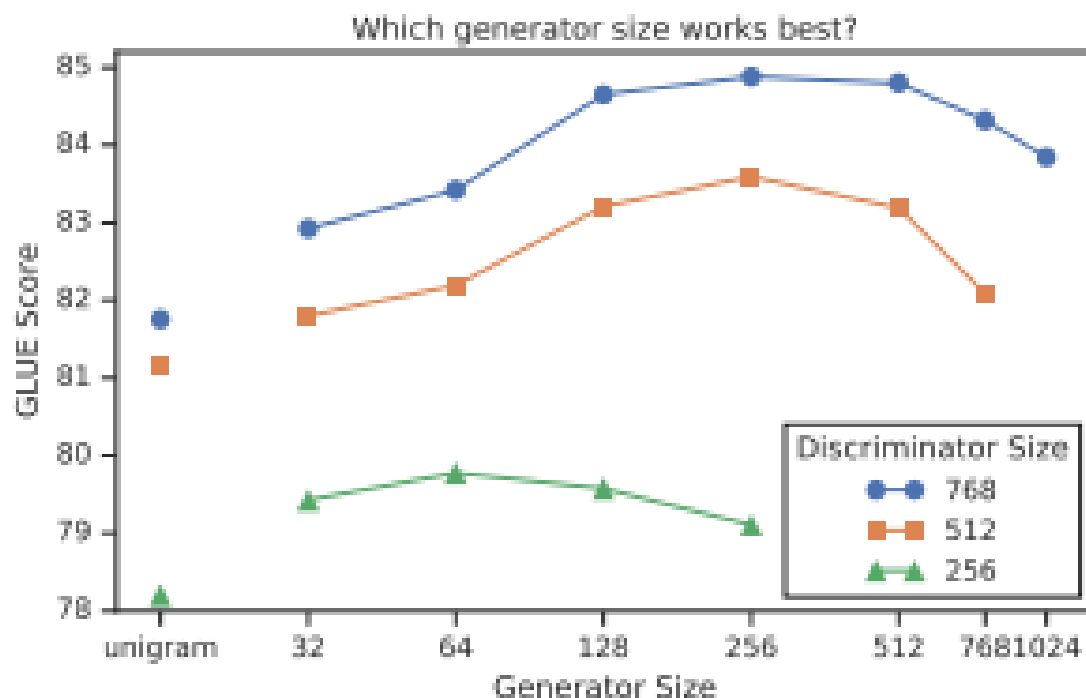
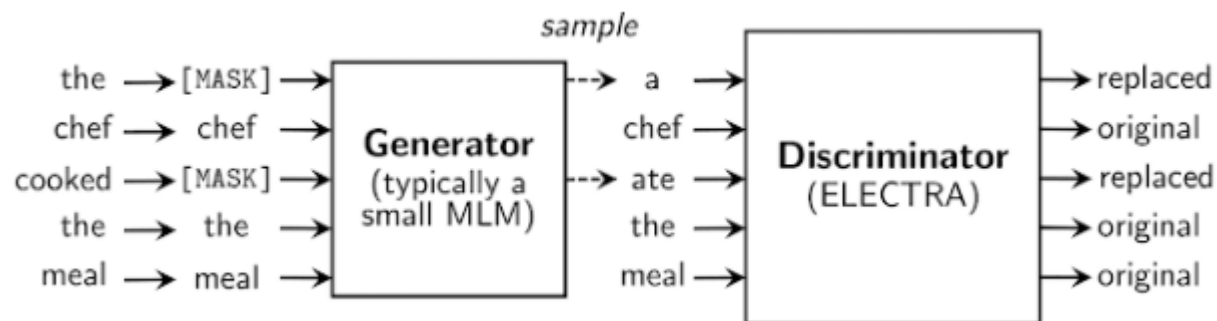
- GLUE - RTE,MNLI,QNLI, MRPC, QQP, STS, SST, and CoLA.
- SQuAD dataset 벤치마크
- BERT와 동일한 사전학습 데이터셋 으로 훈련
(3.3 Billion tokens from Wikipedia and BooksCorpus)
- 모델 구조와 대부분의 하이퍼파라미터는 BERT와 동일
- fine-tuning 을 위해 간단한 linear classifiers 레이어 추가
- fine-tuning시 작은 평가 데이터셋의 랜덤시드 의존성을 고려해
동일 체크포인트에서 평가의 정확도는 10회 중앙값 보고

Experiments : MODEL EXTENSIONS

- Weight Sharing
- Smaller Generators
- Training Algorithms


Experiments : MODEL EXTENSIONS

- Weight Sharing
- Smaller Generators



*FLOPs는 1초당 부동 소수점 연산(곱셈)의 명령 실행 횟수를 나타내는 단위

Training Algorithms



- **Adversarial Contrastive Estimation**


강화 학습을 이용하여 생성 모델을 GAN과 같이 적대적인 방법으로 학습하는 방법

- **Two-stage training**

1. n 단계 동안 L MLM G만 훈련
2. 판별 모델의 weight를 생성 모델의 것으로 초기화하고, 생성모델의 가중치를 고정된 상태로 유지, 판별 모델만 n step 학습 .

*FLOPs는 1초당 부동 소수점 연산(곱셈)의 명령 실행 횟수를 나타내는 단위

Training Algorithms **result**



- **Adversarial Contrastive Estimation**

생성모델에 대해 판별모델의 목적함수로 옮겨 갈때 성능이 증가하였지만, 기존 방식 보다 성능 낮음

- **Two-stage training **problems****

1. poor sample efficiency
2. 생성 모델의 낮은 엔트로피 출력 분포

기존방식 사용

*FLOPs는 1초당 부동 소수점 연산(곱셈)의 명령 실행 횟수를 나타내는 단위


Efficiency Analysis

ELECTRA 15%

원래 ELECTRA와 동일하나, 판별 모델의 loss는 마스킹된 15%의 토큰에서 온 것만을 사용

| Model | ELECTRA | All-Tokens MLM | Replace MLM | ELECTRA 15% | BERT |
|------------|---------|----------------|-------------|-------------|------|
| GLUE score | 85.0 | 84.3 | 82.4 | 82.4 | 82.2 |

Table 5: Compute-efficiency experiments (see text for details).




Efficiency Analysis

Replace MLM

MLM과 비슷하지만, 마스킹된 토큰 [MASK]를 인풋으로 받는 대신, 생성 모델이 만들어낸 토큰으로 대체하여 MLM을 진행

| Model | ELECTRA | All-Tokens MLM | Replace MLM | ELECTRA 15% | BERT |
|------------|---------|----------------|-------------|-------------|------|
| GLUE score | 85.0 | 84.3 | 82.4 | 82.4 | 82.2 |

Table 5: Compute-efficiency experiments (see text for details).




Efficiency Analysis

All-Tokens MLM

BERT와 ELECTRA를 합친 것으로 mask된 토큰이 아닌 모든 토큰을 예측하는 모델이다.

| Model | ELECTRA | All-Tokens MLM | Replace MLM | ELECTRA 15% | BERT |
|------------|---------|----------------|-------------|-------------|------|
| GLUE score | 85.0 | 84.3 | 82.4 | 82.4 | 82.2 |

Table 5: Compute-efficiency experiments (see text for details).



MODEL-DETAIL

Small-model

| Model | Train / Infer FLOPs | Speedup | Params | Train Time + Hardware | GLUE |
|---------------|---------------------|--------------|--------|------------------------|------|
| ELMo | 3.3e18 / 2.6e10 | 19x / 1.2x | 96M | 14d on 3 GTX 1080 GPUs | 71.2 |
| GPT | 4.0e19 / 3.0e10 | 1.6x / 0.97x | 117M | 25d on 8 P6000 GPUs | 78.8 |
| BERT-Small | 1.4e18 / 3.7e9 | 45x / 8x | 14M | 4d on 1 V100 GPU | 75.1 |
| BERT-Base | 6.4e19 / 2.9e10 | 1x / 1x | 110M | 4d on 16 TPUv3s | 82.2 |
| ELECTRA-Small | 1.4e18 / 3.7e9 | 45x / 8x | 14M | 4d on 1 V100 GPU | 79.9 |
| 50% trained | 7.1e17 / 3.7e9 | 90x / 8x | 14M | 2d on 1 V100 GPU | 79.0 |
| 25% trained | 3.6e17 / 3.7e9 | 181x / 8x | 14M | 1d on 1 V100 GPU | 77.7 |
| 12.5% trained | 1.8e17 / 3.7e9 | 361x / 8x | 14M | 12h on 1 V100 GPU | 76.0 |
| 6.25% trained | 8.9e16 / 3.7e9 | 722x / 8x | 14M | 6h on 1 V100 GPU | 74.1 |
| ELECTRA-Base | 6.4e19 / 2.9e10 | 1x / 1x | 110M | 4d on 16 TPUv3s | 85.1 |

ELECTRA-small 결과

같은 계산량의 BERT-small보다 성능이 좋음,
훨씬 계산이 많이 필요했던 GPT보다도 좋음

MODEL-DETAIL

Large Models

| Model | Train FLOPs | Params | CoLA | SST | MRPC | STS | QQP | MNLI | QNLI | RTE | Avg. |
|---------------|----------------|--------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|
| BERT | 1.9e20 (0.27x) | 335M | 60.6 | 93.2 | 88.0 | 90.0 | 91.3 | 86.6 | 92.3 | 70.4 | 84.0 |
| RoBERTa-100K | 6.4e20 (0.90x) | 356M | 66.1 | 95.6 | 91.4 | 92.2 | 92.0 | 89.3 | 94.0 | 82.7 | 87.9 |
| RoBERTa-500K | 3.2e21 (4.5x) | 356M | 68.0 | 96.4 | 90.9 | 92.1 | 92.2 | 90.2 | 94.7 | 86.6 | 88.9 |
| XLNet | 3.9e21 (5.4x) | 360M | 69.0 | 97.0 | 90.8 | 92.2 | 92.3 | 90.8 | 94.9 | 85.9 | 89.1 |
| BERT (ours) | 7.1e20 (1x) | 335M | 67.0 | 95.9 | 89.1 | 91.2 | 91.5 | 89.6 | 93.5 | 79.5 | 87.2 |
| ELECTRA-400K | 7.1e20 (1x) | 335M | 69.3 | 96.0 | 90.6 | 92.1 | 92.4 | 90.5 | 94.5 | 86.8 | 89.0 |
| ELECTRA-1.75M | 3.1e21 (4.4x) | 335M | 69.1 | 96.9 | 90.8 | 92.6 | 92.4 | 90.9 | 95.0 | 88.0 | 89.5 |

<GLUE dev set에 대한 결과>

| Model | Train FLOPs | CoLA | SST | MRPC | STS | QQP | MNLI | QNLI | RTE | WNLI | Avg.* | Score |
|---------|----------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|
| BERT | 1.9e20 (0.06x) | 60.5 | 94.9 | 85.4 | 86.5 | 89.3 | 86.7 | 92.7 | 70.1 | 65.1 | 79.8 | 80.5 |
| RoBERTa | 3.2e21 (1.02x) | 67.8 | 96.7 | 89.8 | 91.9 | 90.2 | 90.8 | 95.4 | 88.2 | 89.0 | 88.1 | 88.1 |
| ALBERT | 3.1e22 (10x) | 69.1 | 97.1 | 91.2 | 92.0 | 90.5 | 91.3 | — | 89.2 | 91.8 | 89.0 | — |
| XLNet | 3.9e21 (1.26x) | 70.2 | 97.1 | 90.5 | 92.6 | 90.4 | 90.9 | — | 88.5 | 92.5 | 89.1 | — |
| ELECTRA | 3.1e21 (1x) | 71.7 | 97.1 | 90.7 | 92.5 | 90.8 | 91.3 | 95.8 | 89.8 | 92.5 | 89.5 | 89.4 |

큰 모델을 학습시켜 보았을 때,
ELECTRA는 XLNet이나 RoBERTa 사전학습에
필요한 계산량의 1/4만 사용해도 비슷한 성능 결과

PRE-TRAINING, FINE-TUNING

PRE-TRAINING DETAILS

- 대부분 BERT와 동일한 하이퍼파라미터를 사용
- 손실에서 판별자 목적의 가중치인 λ 를 50.8로 설정
- 전처리 대신에 즉석에서 결정된 마스킹된 위치와 함께 동적 토큰 마스킹을 사용
- 원본 BERT 논문에서 제안한 Next-sentence-prediction Objective를 사용하지 않음(최근 연구에 따르면 점수가 향상되지 않는 것으로 나타남)
- ELECTRA-Large 모델의 경우 더 높은 마스크 퍼센트 (15 대신 25)를 사용 : Generator가 15% 마스킹으로 높은 정확도를 달성하여 교체된 토큰이 거의 없음을 확인함
- 초기 실험에서 $[1e-4, 2e-4, 3e-4, 5e-4]$ 중 Base 및 Small 모델에 대한 최상의 학습률을 검색하고 $[1, 10, 20, 50, 100]$ 중 λ 를 선택

FINE-TUNING DETAILS

- 라지모델은 대부분 Clark et al.의 하이퍼파라미터를 사용.
- RoBERTa가 더 많은 training epoch(3이 아닌 10까지)를 사용한다는 사실을 알게 된 후, 각 작업에 대해 $[10, 3]$ 중에서 가장 좋은 수의 train epoch를 찾음
- SQuAD의 경우 BERT 및 RoBERTa와 일치하도록 학습 epochs 수를 2로 줄임
- Base모델은 $[3e-5, 5e-5, 1e-4, 1.5e-4]$ 에서 학습률 선택, layer-wise learning-rate decay의 경우 $[0.9, 0.8, 0.7]$ 에서 선택, 나머지는 라지모델과 동일한 하이퍼파라미터를 사용
- 이에 반해 BERT, XLNet, RoBERTa 등의 GLUE에 대한 선행 연구에서 태스크별로 최적의 하이퍼파라미터를 별도로 선택 (동일한 종류의 추가 초매개변수 검색을 수행하면 결과가 약간 향상될 것으로 예상)

Conclusion

- ✓ language representation learning을 위한 새로운 Self-supervised learning task인 **Replaced Token Detection(RTD)**을 제안
- ✓ 핵심은 text encoder가 작은 generator가 만들어낸 고품질의 negative sample과 입력 토큰을 구별하도록 텍스트 인코더를 학습시키는 것
- ✓ MLM과 비교해, 논문의 pre-training objective 계산이 훨씬 효율적이며 downstream task에서 더 좋은 성능을 낼 수 있다는 것을 많은 실험을 통해 확인
- ✓ 이 연구를 통해 컴퓨팅 리소스에 대한 접근 권한이 적은 연구원과 실무자가 적은 컴퓨팅 리소스로도 pre-trained text encoder에 대해 많은 연구개발을 할 수 있길 바람
- ✓ pre-training과 관련된 향후 연구가 절대적 성능 지표만큼 계산량과 파라미터 수 등의 효율성도 함께 고려했으면 하는 바람을 나타냄

Source Code

Source Code

Google-research / electra

<https://github.com/google-research/electra>

ELECTRA

Introduction

ELECTRA is a method for self-supervised language representation learning. It can be used to pre-train transformer networks using relatively little compute. ELECTRA models are trained to distinguish "real" input tokens vs "fake" input tokens generated by another neural network, similar to the discriminator of a [GAN](#). At small scale, ELECTRA achieves strong results even when trained on a single GPU. At large scale, ELECTRA achieves state-of-the-art results on the [SQuAD 2.0](#) dataset.

For a detailed description and experimental results, please refer to our ICLR 2020 paper [ELECTRA: Pre-training Text Encoders as Discriminators Rather Than Generators](#).

This repository contains code to pre-train ELECTRA, including small ELECTRA models on a single GPU. It also supports fine-tuning ELECTRA on downstream tasks including classification tasks (e.g., [GLUE](#)), QA tasks (e.g., [SQuAD](#)), and sequence tagging tasks (e.g., [text chunking](#)).

This repository also contains code for **Electric**, a version of ELECTRA inspired by [energy-based models](#). Electric provides a more principled view of ELECTRA as a "negative sampling" [cloze model](#). It can also efficiently produce [pseudo-likelihood scores](#) for text, which can be used to re-rank the outputs of speech recognition or machine translation systems. For details on Electric, please refer to our EMNLP 2020 paper [Pre-Training Transformers as Energy-Based Cloze Models](#).



clarkkev Merge pull request #88 from PI



finetune



model



pretrain



util



CONTRIBUTING.md



LICENSE



README.md



build_openwebtext_pretraining_datas...



build_pretraining_dataset.py



configure_finetuning.py



configure_pretraining.py



flops_computation.py



run_finetuning.py



run_pretraining.py

Config – configure_pretraining.py

```
# model settings
self.model_size = "small" # one of "small", "base", or "large"
# override the default transformer hparams for the provided model size; see
# modeling.BertConfig for the possible hparams and util.training_utils for
# the defaults
self.model_hparam_overrides = (
    kwargs["model_hparam_overrides"]
    if "model_hparam_overrides" in kwargs else {}
)
self.embedding_size = None # bert hidden size by default
self.vocab_size = 30522 # number of tokens in the vocabulary
self.do_lower_case = True # lowercase the input?
```

```
# batch sizes
self.max_seq_length = 128
self.train_batch_size = 128
self.eval_batch_size = 128
```

```
# model settings
self.model_size = "small" # one of "small", "base", or "large"
# override the default transformer hparams for the provided model size; see
# modeling.BertConfig for the possible hparams and util.training_utils for
# the defaults
self.model_hparam_overrides = (
    kwargs["model_hparam_overrides"]
    if "model_hparam_overrides" in kwargs else {}
)
self.embedding_size = None # bert hidden size by default
self.vocab_size = 30522 # number of tokens in the vocabulary
self.do_lower_case = True # lowercase the input?
```

```
# generator settings
self.uniform_generator = False # generator is uniform at random
self.two_tower_generator = False # generator is a two-tower cloze model
self.untied_generator_embeddings = False # tie generator/discriminator
# token embeddings?
self.untied_generator = True # tie all generator/discriminator weights?
self.generator_layers = 1.0 # frac of discriminator layers for generator
self.generator_hidden_size = 0.25 # frac of discrim hidden size for gen
self.disallow_correct = False # force the generator to sample incorrect
# tokens (so 15% of tokens are always
# fake)
self.temperature = 1.0 # temperature for sampling from generator
```

```
def get_bert_config(config):
    """Get model hyperparameters based on a pretraining/finetuning config"""
    if config.model_size == "large":
        args = {"hidden_size": 1024, "num_hidden_layers": 24}
    elif config.model_size == "base":
        args = {"hidden_size": 768, "num_hidden_layers": 12}
    elif config.model_size == "small":
        args = {"hidden_size": 256, "num_hidden_layers": 12}
    else:
        raise ValueError("Unknown model size", config.model_size)
    args["vocab_size"] = config.vocab_size
    args.update(**config.model_hparam_overrides)
    # by default the ff size and num attn heads are determined by the hidden size
    args["num_attention_heads"] = max(1, args["hidden_size"] // 64)
    args["intermediate_size"] = 4 * args["hidden_size"]
    args.update(**config.model_hparam_overrides)
    return modeling.BertConfig.from_dict(args)
```


Tokenizer

road
wasn
although
due
major
died
village
third
knew
2016
asked
turned
st
wanted
say
##p
together
received
main
son
served
different
##en
behind
himself
felt
members
power
football
law
voice
play
##in
near
park
history
30
having
2005

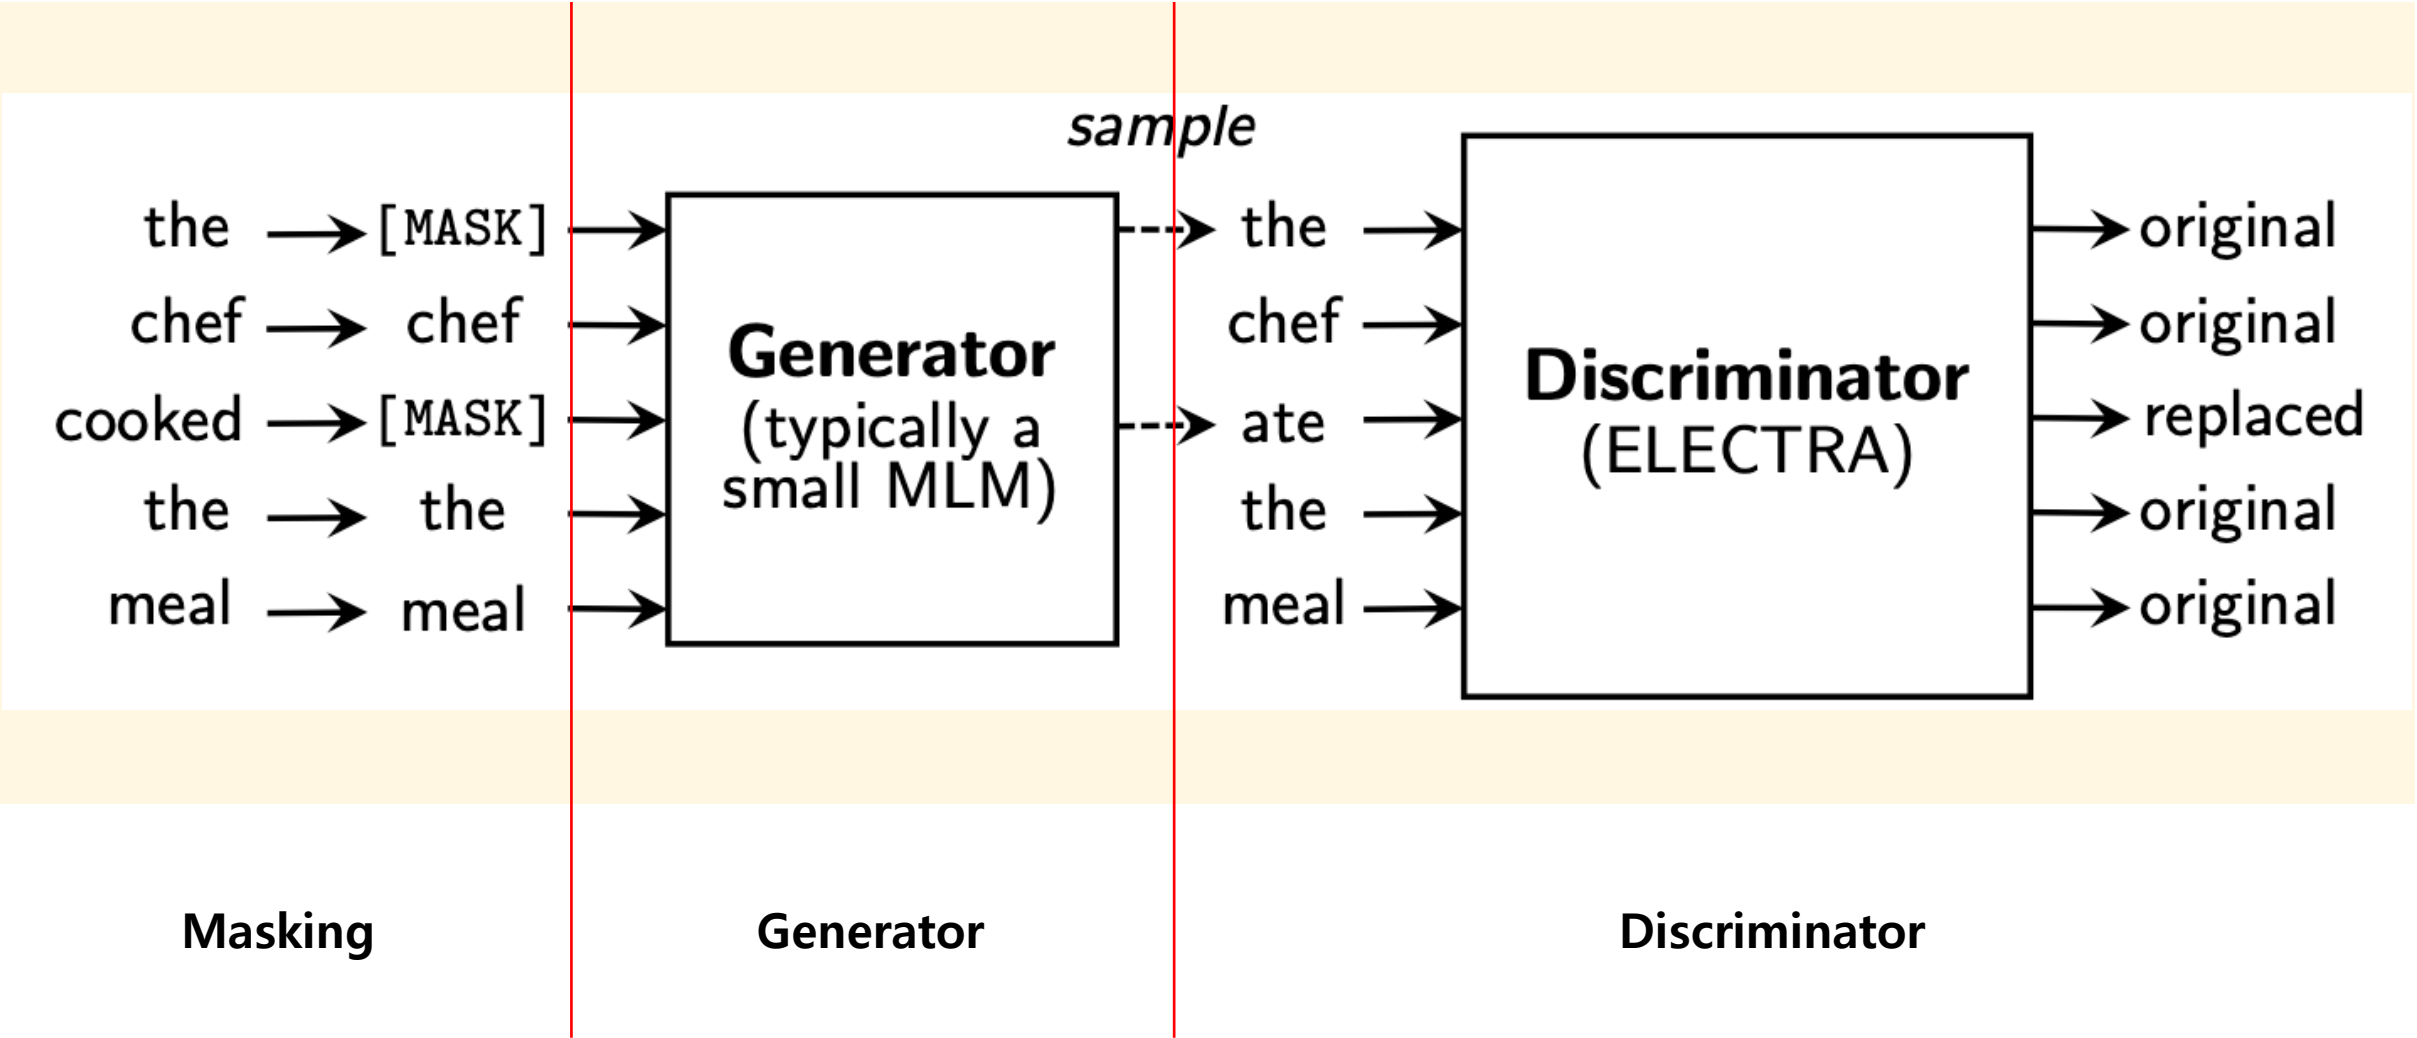
[Vocab]

Wordpiece Vocabulary

1929, "空": 1930, "立": 1931, "章": 1932, "竹": 1933, "系": 1934, "美": 1935, "義": 1936, "耳": 1937, "良": 1938, "十": 1939, "花": 1940, "英": 1941, "華": 1942, "葉": 1943, "藤": 1944, "行": 1945, "街": 1946, "郎": 1958, "郡": 1959, "部": 1960, "都": 1961, "里": 1962, "野": 1963, "金": 1964, "鈴": 1965, "鎮": 1966, "長": 1967, "門": 1968, "間": 1969, "β": 1970, "阿": 1971, "陳": 1972, "陽": 1973, "雄": 1974, "青": 1975, " (": 1987, ")": 1988, ",": 1989, "一": 1990, ".": 1991, "/" : 1992, " ": 1993, " ? ": 1994, " ~ ": 1995, "the": 1996, "of": 1997, "and": 1998, "in": 1999, "to": 2000, "was": 2001, "he": 2002, "is": 2003, "": 2004, "##s": 2015, "she": 2016, "you": 2017, "had": 2018, "an": 2019, "were": 2020, "but": 2021, "be": 2022, "this": 2023, "are": 2024, "not": 2025, "my": 2026, "they": 2027, "one": 2028, "which": 2029, "who": 2030, "out": 2041, "been": 2042, "when": 2043, "after": 2044, "there": 2045, "into": 2046, "new": 2047, "two": 2048, "its": 2049, "##a": 2050, "time": 2051, "would": 2052, "no": 2053, "what": 2054, "if": 2065, "I like": 2066, "back": 2067, "them": 2068, "only": 2069, "some": 2070, "could": 2071, "##i": 2072, "where": 2073, "just": 2074, "##ing": 2075, "during": 2076, "before": 2077, "##n": 2078, "world": 2088, "may": 2089, "between": 2090, "down": 2091, "well": 2092, "three": 2093, "##d": 2094, "year": 2095, "while": 2096, "will": 2097, "##ed": 2098, "##r": 2099, "##y": 2100, "later": 2101, "people": 2111, "part": 2112, "know": 2113, "against": 2114, "your": 2115, "many": 2116, "second": 2117, "university": 2118, "both": 2119, "national": 2120, "##er": 2121, "these": 2122, "don": 2123, "knew": 2134, "##ly": 2135, "team": 2136, "american": 2137, "because": 2138, "de": 2139, "##l": 2140, "born": 2141, "united": 2142, "film": 2143, "since": 2144, "still": 2145, "long": 2146, "work": 2147, "right": 2157, "man": 2158, "eyes": 2159, "house": 2160, "season": 2161, "war": 2162, "states": 2163, "including": 2164, "took": 2165, "life": 2166, "north": 2167, "same": 2168, "each": 2169, "called": 2170, "won": 2179, "area": 2180, "here": 2181, "going": 2183, "10": 2184, "away": 2185, "series": 2186, "left": 2187, "home": 2188, "music": 2189, "best": 2190, "make": 2191, "hand": 2192, "number": 2193, "end": 2202, "good": 2203, "too": 2204, "following": 2206, "released": 2207, "game": 2208, "played": 2209, "little": 2210, "began": 2211, "district": 2212, "##m": 2213, "old": 2214, "want": 2215, "west": 2225, "##u": 2226, "face": 2227, "think": 2228, "##es": 2229, "2010": 2230, "government": 2231, "##h": 2232, "march": 2233, "came": 2234, "small": 2235, "general": 2236, "town": 2237, "june": 2246, "along": 2247, "international": 2248, "2011": 2249, "air": 2250, "july": 2251, "club": 2252, "went": 2253, "january": 2254, "october": 2255, "our": 2256, "august": 2257, "april": 2258, "york": 2259, "father": 2269, "public": 2270, "##us": 2271, "come": 2272, "men": 2273, "five": 2274, "set": 2275, "station": 2276, "church": 2277, "##c": 2278, "next": 2279, "former": 2280, "november": 2281, "room": 2291, "let": 2292, "love": 2293, "2006": 2294, "though": 2295, "every": 2296, "2014": 2297, "look": 2298, "song": 2299, "water": 2300, "century": 2301, "without": 2302, "body": 2303, "black": 2304, "ulation": 2313, "river": 2314, "named": 2315, "band": 2316, "white": 2317, "started": 2318, "##an": 2319, "once": 2320, "15": 2321, "20": 2322, "should": 2323, "18": 2324, "2015": 2325, "service": 2326, "children": 2336, "february": 2337, "book": 2338, "why": 2339, "11": 2340, "door": 2341, "need": 2342, "president": 2343, "order": 2344, "final": 2345, "road": 2346, "wasn": 2347, "although": 2348, "st": 2358, "wanted": 2359, "say": 2360, "##p": 2361, "together": 2362, "received": 2363, "main": 2364, "son": 2365, "served": 2366, "different": 2367, "##en": 2368, "behind": 2369, "himself": 2370, "r": 2379, "park": 2380, "history": 2381, "30": 2382, "having": 2383, "2005": 2384, "16": 2385, "##man": 2386, "saw": 2387, "mother": 2388, "##al": 2389, "army": 2390, "point": 2391, "front": 2392, "help": 2402, "14": 2403, "put": 2404, "published": 2405, "country": 2406, "division": 2407, "across": 2408, "told": 2409, "13": 2410, "often": 2411, "ever": 2412, "french": 2413, "london": 2414, "cent": 2424, "tell": 2425, "among": 2426, "species": 2427, "really": 2428, "according": 2429, "central": 2430, "half": 2431, "2004": 2432, "form": 2433, "original": 2434, "gave": 2435, "office": 2436, "make": 2446, "german": 2446, "player": 2447, "run": 2448, "business": 2449, "woman": 2450, "community": 2451, "cup": 2452, "might": 2453, "million": 2454, "land": 2455, "2000": 2456, "court": 2457, "development": 2467, "become": 2468, "sure": 2469, "research": 2470, "almost": 2471, "director": 2472, "council": 2473, "la": 2474, "##2": 2475, "career": 2476, "things": 2477, "using": 2478, "island": 2479, "##z": 2480, "support": 2490, "control": 2491, "field": 2492, "students": 2493, "2003": 2494, "education": 2495, "married": 2496, "##b": 2497, "nothing": 2498, "worked": 2499, "others": 2500, "record": 2501, "try": 2510, "established": 2511, "non": 2512, "returned": 2513, "feel": 2514, "does": 2515, "title": 2516, "written": 2517, "thing": 2518, "feet": 2519, "william": 2520, "far": 2521, "co": 2522, "assoc": 2531, "##na": 2532, "department": 2533, "hall": 2534, "role": 2535, "various": 2536, "production": 2537, "21": 2538, "19": 2539, "heart": 2540, "2001": 2541, "living": 2542, "fire": 2543, "vers": 2552, "case": 2553, "society": 2554, "region": 2555, "present": 2556, "radio": 2557, "period": 2558, "looking": 2559, "least": 2560, "total": 2561, "keep": 2562, "england": 2563, "wife": 2564, "pro": 2574, "##6": 2575, "political": 2576, "george": 2577, "services": 2578, "taken": 2579, "created": 2580, "##7": 2581, "further": 2582, "table": 2583, "reached": 2584, "david": 2585, "union": 2586, "##x": 2595, "appeared": 2596, "position": 2597, "ground": 2598, "lead": 2599, "rock": 2600, "dark": 2601, "election": 2602, "23": 2603, "board": 2604, "france": 2605, "hair": 2606, "course": 2607, "moment": 2617, "##te": 2618, "someone": 2619, "##8": 2620, "summer": 2621, "project": 2622, "announced": 2623, "san": 2624, "less": 2625, "wrote": 2626, "past": 2627, "followed": 2628, "##5": 2629, "ne": 2638, "1999": 2639, "design": 2640, "considered": 2641, "northern": 2642, "god": 2643, "stop": 2644, "battle": 2645, "toward": 2646, "european": 2647, "outside": 2648, "described": 2649, "track

[Token]

ELECTRA Structure



Pretrain-Model - masking

```
# Mask the input
unmasked_inputs = pretrain_data.features_to_inputs(features)
masked_inputs = pretrain_helpers.mask(
    config, unmasked_inputs, config.mask_prob)
```

[run_pretraining.py]

```
96 # model inputs - it's a bit nicer to use a namedtuple rather than keep the
97 # features as a dict
98 Inputs = collections.namedtuple(
99     "Inputs", ["input_ids", "input_mask", "segment_ids", "masked_lm_positions",
100               "masked_lm_ids", "masked_lm_weights"])
```

[pretrain_data.py]

```
# Update the input ids
replace_with_mask_positions = masked_lm_positions * tf.cast(
    tf.less(tf.random.uniform([B, N]), 0.85), tf.int32)
inputs_ids, _ = scatter_update(
    inputs.input_ids, tf.fill([B, N], vocab["[MASK]"]),
    replace_with_mask_positions)

return pretrain_data.get_updated_inputs(
    inputs,
    input_ids=tf.stop_gradient(inputs_ids),
    masked_lm_positions=masked_lm_positions,
    masked_lm_ids=masked_lm_ids,
    masked_lm_weights=masked_lm_weights
)
```

[pretrain_helpers.py]

$$m_i \sim \text{unif}\{1, n\} \text{ for } i = 1 \text{ to } k$$

[마스킹 위치 랜덤]

$$x^{\text{masked}} = \text{REPLACE}(x, m, [\text{MASK}])$$

[MASK 토큰으로 대체]

B – Batch Size

L – Sequence Length

D – Depth

N – modeling.get_shape_list(position)[1]

Pretrain-Model – Generator - Config

```
# generator settings
self.uniform_generator = False # generator is uniform at random
self.two_tower_generator = False # generator is a two-tower cloze model
self.untied_generator_embeddings = False # tie generator/discriminator
                                         # token embeddings?

self.untied_generator = True # tie all generator/discriminator weights?
self.generator_layers = 1.0 # frac of discriminator layers for generator
self.generator_hidden_size = 0.25 # frac of discrim hidden size for gen
self.disallow_correct = False # force the generator to sample incorrect
                               # tokens (so 15% of tokens are always
                               # fake)

self.temperature = 1.0 # temperature for sampling from generator
```

[Generator setting]

[Generator Config List]

1. Generator is uniform at random
2. Generator is a two-tower cloze model
3. Tie generator / discriminator token embedding
4. Tie all generator / discriminator weights

Pretrain-Model – Generator1

```
61     if config.uniform_generator:
62         # simple generator sampling fakes uniformly at random
63         mlm_output = self._get_masked_lm_output(masked_inputs, None)

164     def _get_masked_lm_output(self, inputs: pretrain_data.Inputs, model):
165         """Masked language modeling softmax layer."""
166         with tf.variable_scope("generator_predictions"):
167             if self._config.uniform_generator:
168                 logits = tf.zeros(self._bert_config.vocab_size)
169                 logits_tiled = tf.zeros(
170                     modeling.get_shape_list(inputs.masked_lm_ids) +
171                     [self._bert_config.vocab_size])
172                 logits_tiled += tf.reshape(logits, [1, 1, self._bert_config.vocab_size])
173                 logits = logits_tiled
174             else:
175                 relevant_reprs = pretrain_helpers.gather_positions(
176                     model.get_sequence_output(), inputs.masked_lm_positions)
177                 logits = get_token_logits(
178                     relevant_reprs, model.get_embedding_table(), self._bert_config)
179             return get_softmax_output(
180                 logits, inputs.masked_lm_ids, inputs.masked_lm_weights,
181                 self._bert_config.vocab_size)
```

$$p_G(x_t | x^{masked}) = \exp(e(x_t)^T h_G(x^{masked})_t) / \sum_{x'} \exp(e(x')^T h_G(x^{masked})_t)$$

$$p_G(x_t | x^{masked})$$

⇒ x^{masked} 라는 시퀀스 주어졌을 때 특정 위치 t 에 단어 x_t 가 들어갈 확률 $\exp(e \bullet x_i) / \sum_j \exp(e \bullet x_j)$

⇒ softmax 함수

$$e(x_t)^T h_G(x^{masked})_t$$

⇒ token을 임베딩(e)하여 example의 hidden state 값을 곱

```
271     def get_softmax_output(logits, targets, weights, vocab_size):
272         oh_labels = tf.one_hot(targets, depth=vocab_size, dtype=tf.float32)
273         preds = tf.argmax(logits, axis=-1, output_type=tf.int32)
274         probs = tf.nn.softmax(logits)
275         log_probs = tf.nn.log_softmax(logits)
276         label_log_probs = -tf.reduce_sum(log_probs * oh_labels, axis=-1)
277         numerator = tf.reduce_sum(weights * label_log_probs)
278         denominator = tf.reduce_sum(weights) + 1e-6
279         loss = numerator / denominator
280         SoftmaxOutput = collections.namedtuple(
281             "SoftmaxOutput", ["logits", "probs", "loss", "per_example_loss", "preds",
282                               "weights"])
283         return SoftmaxOutput(
284             logits=logits, probs=probs, per_example_loss=label_log_probs,
285             loss=loss, preds=preds, weights=weights)
```

[run_pretraining.py]

Pretrain-Model – Generator2

```
class TwoTowerClozeTransformer(object):
    """Build a two-tower Transformer used as Electric's generator."""

    def __init__(self, config, bert_config, inputs: pretrain_data.Inputs,
                 is_training, embedding_size):
        ltr = build_transformer(
            config, inputs, is_training, bert_config,
            untied_embeddings=config.untied_generator_embeddings,
            embedding_size=(None if config.untied_generator_embeddings
                           else embedding_size),
            scope="generator_ltr", ltr=True)
        rtl = build_transformer(
            config, inputs, is_training, bert_config,
            untied_embeddings=config.untied_generator_embeddings,
            embedding_size=(None if config.untied_generator_embeddings
                           else embedding_size),
            scope="generator_rtl", rtl=True)

        ltr_reprs = ltr.get_sequence_output()
        rtl_reprs = rtl.get_sequence_output()
        self._sequence_output = tf.concat([roll(ltr_reprs, -1),
                                             roll(rtl_reprs, 1)], -1)

        self._embedding_table = ltr.embedding_table

    def get_sequence_output(self):
        return self._sequence_output

    def get_embedding_table(self):
        return self._embedding_table
```




Figure 1: Illustration of the model. Block_i is a standard transformer decoder block. Green blocks operate left to right by masking future time-steps and blue blocks operate right to left. At the top, states are combined with a standard multi-head self-attention module whose output is fed to a classifier that predicts the center token.




Figure 2: Illustration of fine-tuning for a single-sentence task where the output of the first and last token is fed to a task-specific classifier (W). Masking for the final combination layer (comb) is removed which results in representations based on all forward and backward states (cf. Figure 1).

Cloze-driven Pretraining of Self-attention Networks

Pretrain-Model – Generator3

```
else:
    # small masked language model generator
    generator = build_transformer(
        config, masked_inputs, is_training, generator_config,
        embedding_size=(None if config.untied_generator_embeddings
                        else embedding_size),
        untied_embeddings=config.untied_generator_embeddings,
        scope="generator")
    mlm_output = self._get_masked_lm_output(masked_inputs, generator)
```

```
338 def get_generator_config(config: configure_pretraining.PretrainingConfig,
339                          bert_config: modeling.BertConfig):
340     """Get model config for the generator network."""
341     gen_config = modeling.BertConfig.from_dict(bert_config.to_dict())
342     gen_config.hidden_size = int(round(
343         bert_config.hidden_size * config.generator_hidden_size))
344     gen_config.num_hidden_layers = int(round(
345         bert_config.num_hidden_layers * config.generator_layers))
346     gen_config.intermediate_size = 4 * gen_config.hidden_size
347     gen_config.num_attention_heads = max(1, gen_config.hidden_size // 64)
348     return gen_config
```

```
# generator settings
self.uniform_generator = False # generator is uniform at random
self.two_tower_generator = False # generator is a two-tower cloze model
self.untied_generator_embeddings = False # tie generator/discriminator
                                         # token embeddings?
self.untied_generator = True # tie all generator/discriminator weights?
self.generator_layers = 1.0 # frac of discriminator layers for generator
self.generator_hidden_size = 0.25 # frac of discrim hidden size for gen
self.disallow_correct = False # force the generator to sample incorrect
                               # tokens (so 15% of tokens are always
                               # fake)
self.temperature = 1.0 # temperature for sampling from generator
```

[Generator setting]

[run_pretraining.py]

Pretrain-Model – Generator4

else:

```
# full-sized masked language model generator if using BERT objective or if
# the generator and discriminator have tied weights
```

```
generator = build_transformer(
    config, masked_inputs, is_training, self._bert_config,
    embedding_size=embedding_size)
```

```
mlm_output = self._get_masked_lm_output(masked_inputs, generator)
```

```
96 def get_bert_config(config):
97     """Get model hyperparameters based on a pretraining/finetuning config"""
98     if config.model_size == "large":
99         args = {"hidden_size": 1024, "num_hidden_layers": 24}
100     elif config.model_size == "base":
101         args = {"hidden_size": 768, "num_hidden_layers": 12}
102     elif config.model_size == "small":
103         args = {"hidden_size": 256, "num_hidden_layers": 12}
104     else:
105         raise ValueError("Unknown model size", config.model_size)
106     args["vocab_size"] = config.vocab_size
107     args.update(**config.model_hparam_overrides)
108     # by default the ff size and num attn heads are determined by the hidden size
109     args["num_attention_heads"] = max(1, args["hidden_size"] // 64)
110     args["intermediate_size"] = 4 * args["hidden_size"]
111     args.update(**config.model_hparam_overrides)
112     return modeling.BertConfig.from_dict(args)
```

Pretrain-Model - Generator

Discriminator D input 생성

```
fake_data = self._get_fake_data(masked_inputs, mlm_output.logits)
self.mlm_output = mlm_output
self.total_loss = config.gen_weight * (
    cloze_output.loss if config.two_tower_generator else mlm_output.loss)
```

[Update Generator loss]

[MASK]에서 $p_G(x_t|x)$ 으로 샘플링한 토큰으로 치환(corrupt)

$$x^{corrupt} = REPLACE(x, m, \hat{x})$$

$$\hat{x} \sim p_G(x_i|x^{masked}) \text{ for } i \in m$$

```
FakedData = collections.namedtuple("FakedData", [
    "inputs", "is_fake_tokens", "sampled_tokens"])
```

[FakedData Structure]

Pretrain-Model - Discriminator

```
99     # Discriminator
100    disc_output = None
101    if config.electra_objective or config.electric_objective:
102        discriminator = build_transformer(
103            config, fake_data.inputs, is_training, self._bert_config,
104            reuse=not config.untied_generator, embedding_size=embedding_size)
105    disc_output = self._get_discriminator_output(
106        fake_data.inputs, discriminator, fake_data.is_fake_tokens,
107        cloze_output)
108    self.total_loss += config.disc_weight * disc_output.loss
```

[run_pretraining.py]

Pretrain-Model - Discriminator

```
183 def _get_discriminator_output(  
184     self, inputs, discriminator, labels, cloze_output=None):  
185     """Discriminator binary classifier."""  
186     with tf.variable_scope("discriminator_predictions"):  
187         hidden = tf.layers.dense(  
188             discriminator.get_sequence_output(),  
189             units=self._bert_config.hidden_size,  
190             activation=modeling.get_activation(self._bert_config.hidden_act),  
191             kernel_initializer=modeling.create_initializer(  
192                 self._bert_config.initializer_range))  
193         logits = tf.squeeze(tf.layers.dense(hidden, units=1), -1)  
194         if self._config.electric_objective:  
195             log_q = tf.reduce_sum(  
196                 tf.nn.log_softmax(cloze_output.logits) * tf.one_hot(  
197                     inputs.input_ids, depth=self._bert_config.vocab_size,  
198                     dtype=tf.float32), -1)  
199             log_q = tf.stop_gradient(log_q)  
200             logits += log_q  
201             logits += tf.log(self._config.mask_prob / (1 - self._config.mask_prob))  
202  
203         weights = tf.cast(inputs.input_mask, tf.float32)  
204         labelsf = tf.cast(labels, tf.float32)  
205         losses = tf.nn.sigmoid_cross_entropy_with_logits(  
206             logits=logits, labels=labelsf) * weights  
207         per_example_loss = (tf.reduce_sum(losses, axis=-1) /  
208             (1e-6 + tf.reduce_sum(weights, axis=-1)))  
209         loss = tf.reduce_sum(losses) / (1e-6 + tf.reduce_sum(weights))  
210         probs = tf.nn.sigmoid(logits)  
211         preds = tf.cast(tf.round((tf.sign(logits) + 1) / 2), tf.int32)  
212         DiscOutput = collections.namedtuple(  
213             "DiscOutput", ["loss", "per_example_loss", "probs", "preds",  
214                 "labels"])
```

Target Class (이진)

- Original : 원본 문장의 토큰과 같은 토큰
- Replaced : generator G가 만든 토큰

$$D(x^{corrupt}, t) = \text{sigmoid}(w^T h_D(x^{corrupt})_t)$$

[Discriminator 공식]

Finetune-Model

```
def model_fn(features, labels, mode, params):  
    """The `model_fn` for TPUEstimator."""  
    utils.log("Building model...")  
    is_training = (mode == tf.estimator.ModeKeys.TRAIN)  
    model = FinetuningModel(  
        config, tasks, is_training, features, num_train_steps)
```

```
# Add specific tasks  
self.outputs = {"task_id": features["task_id"]}  
losses = []  
for task in tasks:  
    with tf.variable_scope("task_specific/" + task.name):  
        task_losses, task_outputs = task.get_prediction_module(  
            bert_model, features, is_training, percent_done)  
        losses.append(task_losses)  
        self.outputs[task.name] = task_outputs  
self.loss = tf.reduce_sum(  
    tf.stack(losses, -1) *  
    tf.one_hot(features["task_id"], len(config.task_names)))
```

- Classification_tasks
- Qa_tasks
- Tagging_tasks
-

[run_finetuning.py]

Thank you

Reference

ELECTRA-Code - <https://github.com/google-research/electra>

Vocab - <https://huggingface.co/google/electra-small-generator/resolve/main/vocab.txt>

Tokenizer - <https://huggingface.co/google/electra-small-generator/resolve/main/tokenizer.json>

Paper

ELECTRA - <https://arxiv.org/abs/2003.10555>

BERT - <https://arxiv.org/abs/1810.04805>

Attention - <https://arxiv.org/abs/1706.03762>

Seq2Seq - <https://arxiv.org/abs/1409.3215>

LSTM / RNN

<https://static.googleusercontent.com/media/research.google.com/ko//pubs/archive/43905.pdf>

Cloze-driven Pretraining of Self-attention Networks (Two Tower Cloze Transformer)

<https://arxiv.org/abs/1903.07785>